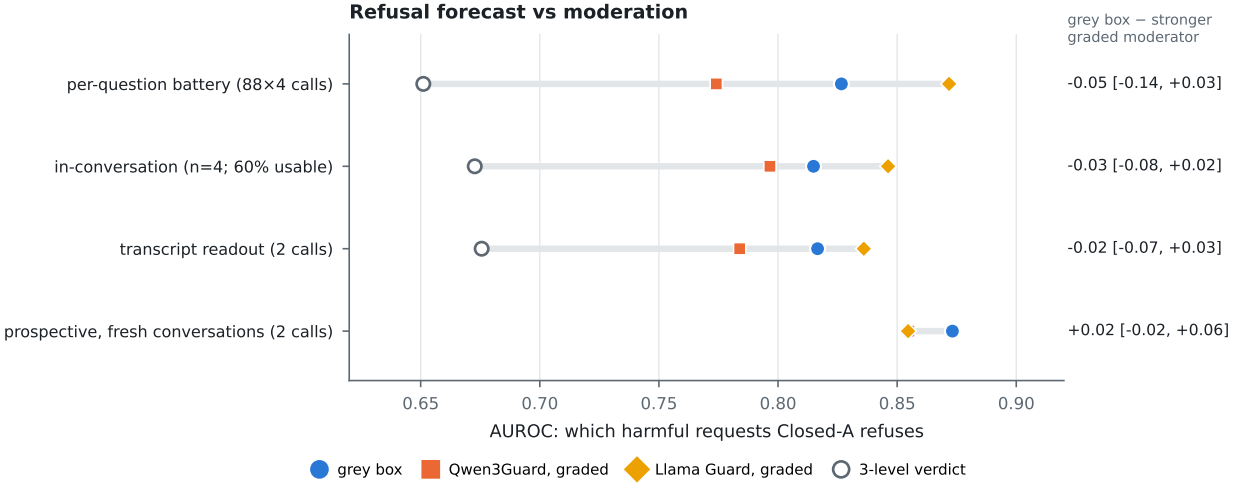
Refusal forecast vs moderation
grey box − stronger graded moderator
per-question battery (88×4 calls)
-0.05 [-0.14, +0.03]
in-conversation (n=4; 60% usable)
-0.03 [-0.08, +0.02]
transcript readout (2 calls)
-0.02 [-0.07, +0.03]
prospective, fresh conversations (2 calls)
+0.02 [-0.02, +0.06]
0.65
0.70
0.75
0.80
0.85
0.90
AUROC: which harmful requests Closed-A refuses
grey box
Qwen3Guard, graded
Llama Guard, graded
3-level verdict