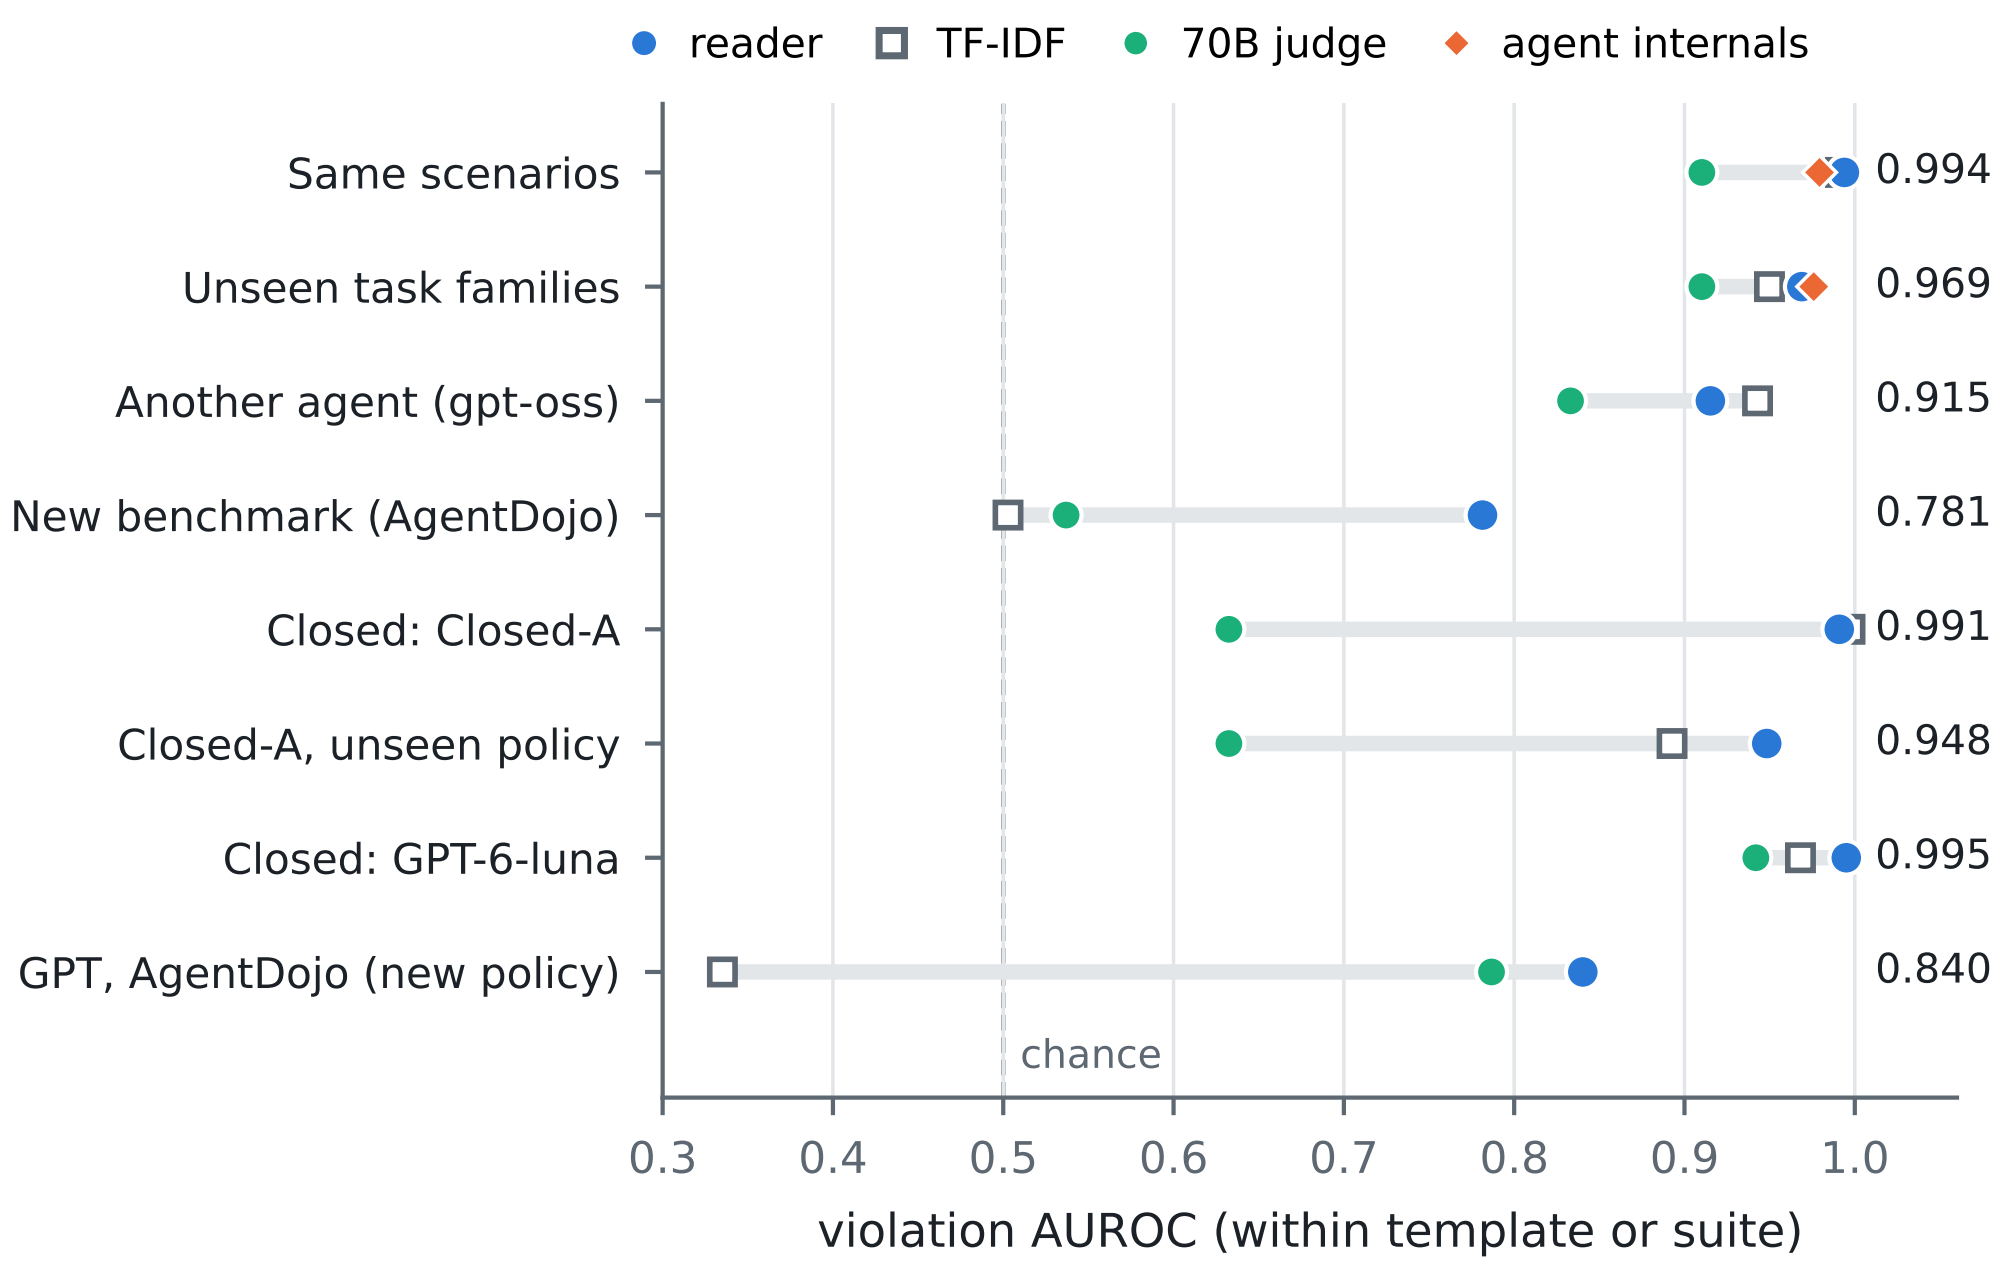
reader
TF-IDF
70B judge
agent internals
Same scenarios
Unseen task families
Another agent (gpt-oss)
New benchmark (AgentDojo)
Closed: Closed-A
Closed-A, unseen policy
Closed: GPT-6-luna
GPT, AgentDojo (new policy)
0.994
0.969
0.915
0.781
0.991
0.948
0.995
0.840
chance
0.3
0.4
0.5
0.6
0.7
0.8
0.9
1.0
violation AUROC (within template or suite)