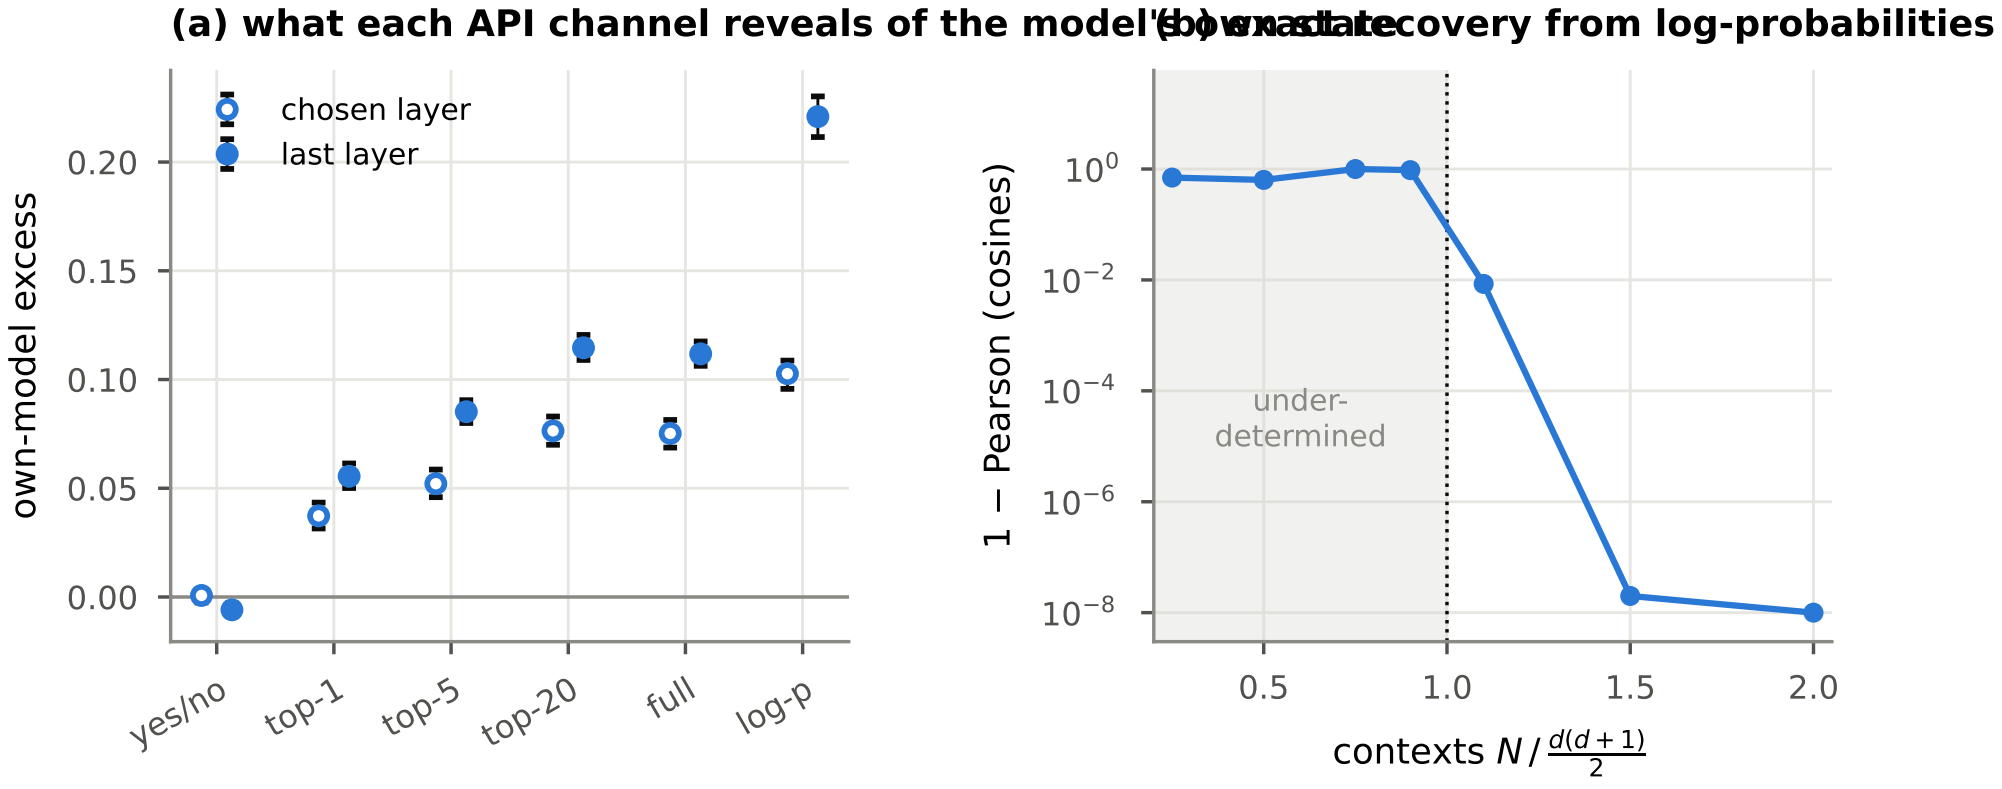
(a) what each API channel reveals of the model's own state
(b) exact recovery from log-probabilities
chosen layer
last layer
own-model excess
0.20
0.15
0.10
0.05
0.00
yes/no
top-1
top-5
top-20
full
log-p
1 − Pearson (cosines)
10^{0}
10^{-2}
10^{-4}
10^{-6}
10^{-8}
under-
determined
0.5
1.0
1.5
2.0
contexts $N / \frac{d(d+1)}{2}$