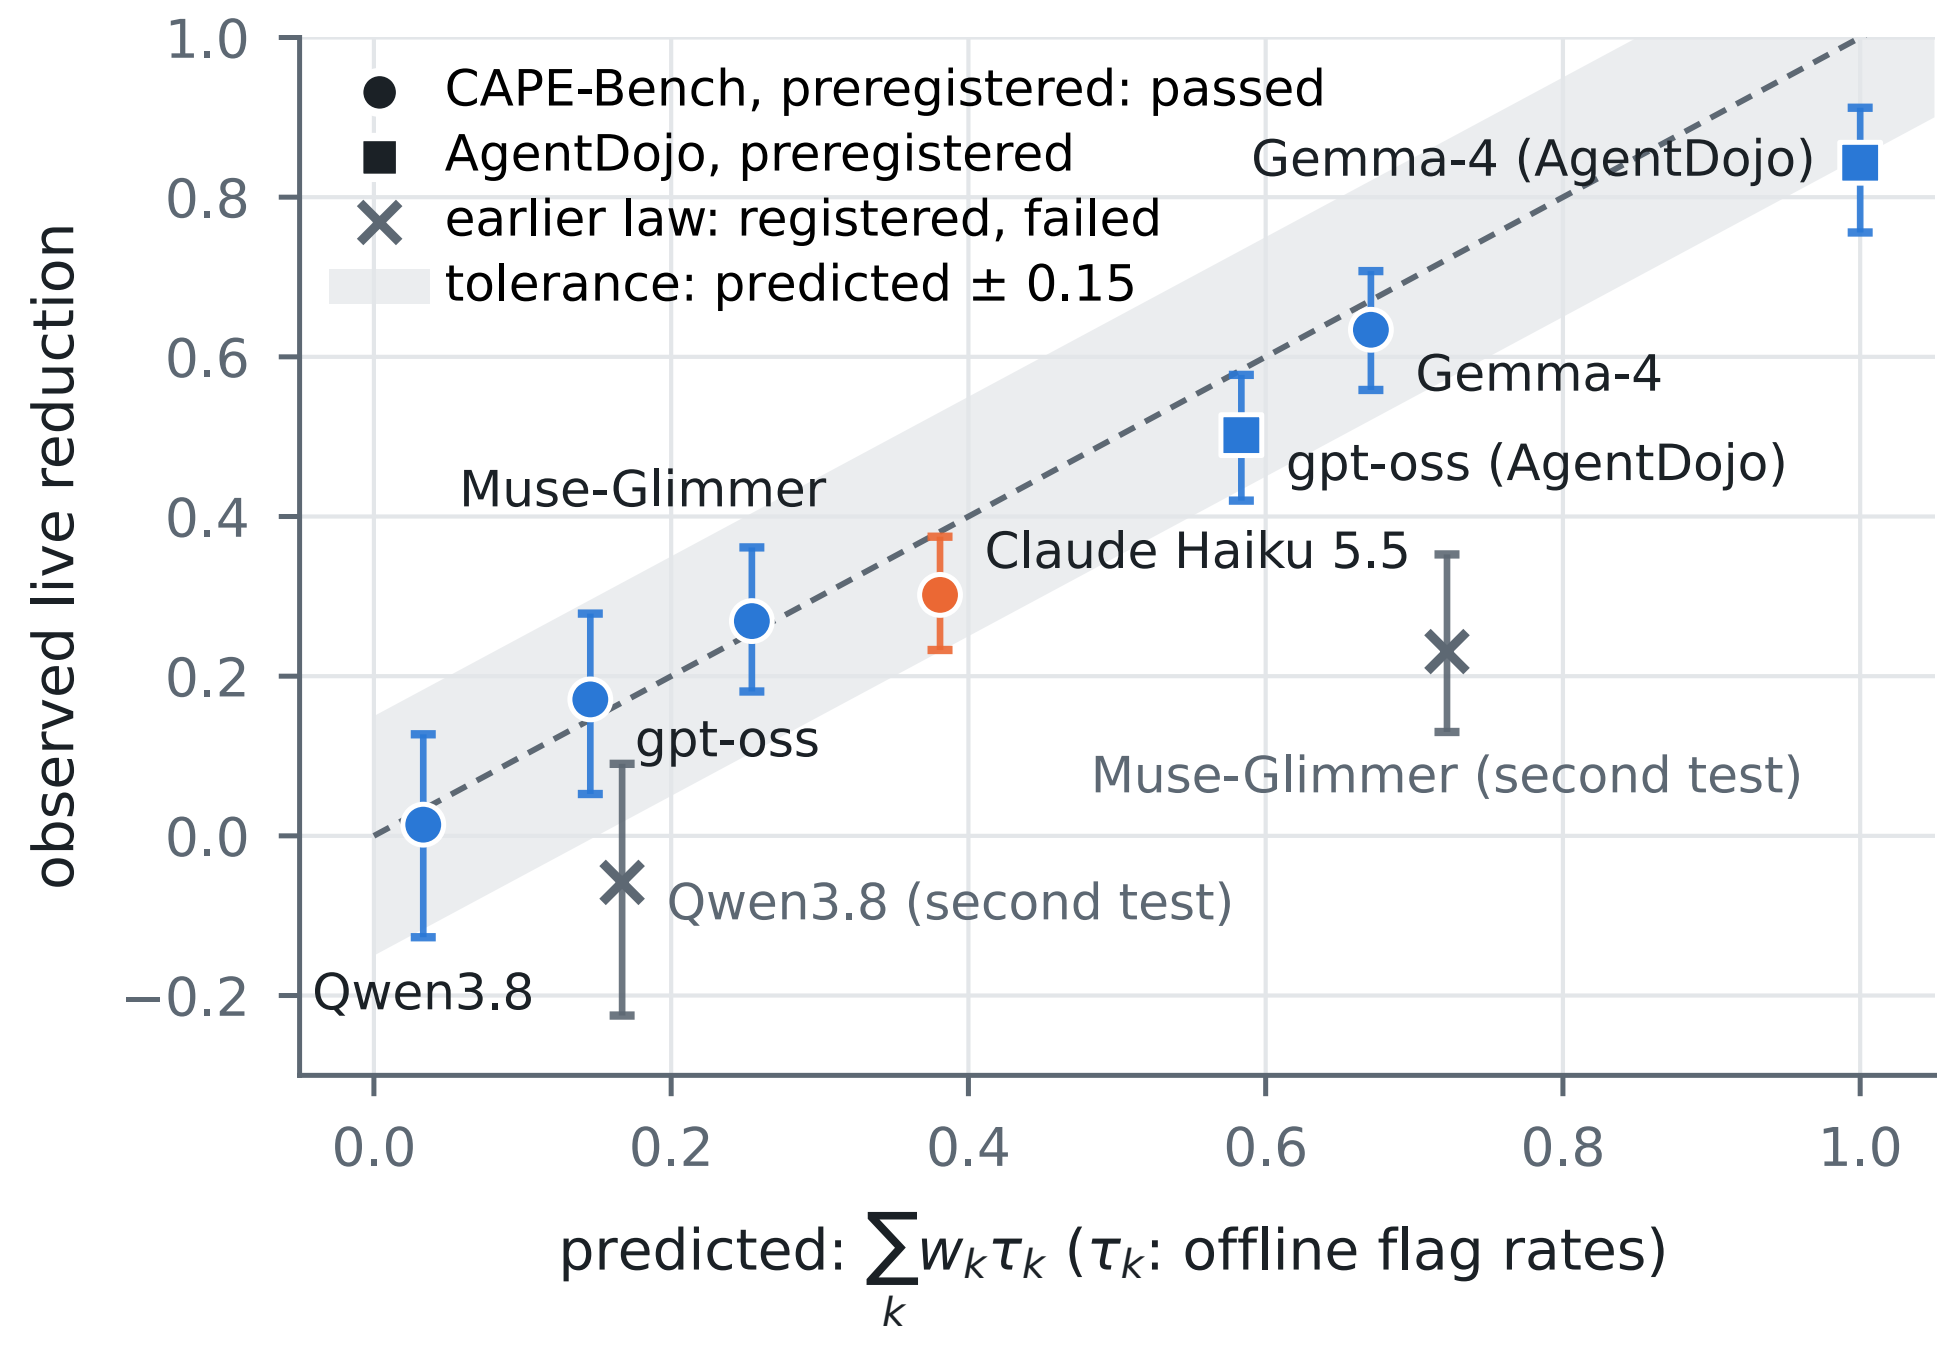
CAPE-Bench, preregistered: passed
AgentDojo, preregistered
earlier law: registered, failed
tolerance: predicted ± 0.15
Gemma-4 (AgentDojo)
Gemma-4
gpt-oss (AgentDojo)
Muse-Glimmer
Claude Haiku 5.5
gpt-oss
Muse-Glimmer (second test)
Qwen3.8 (second test)
Qwen3.8
observed live reduction
predicted: $\sum_k w_k \tau_k$ (τ_k: offline flag rates)
1.0
0.8
0.6
0.4
0.2
0.0
−0.2
0.0
0.2
0.4
0.6
0.8
1.0