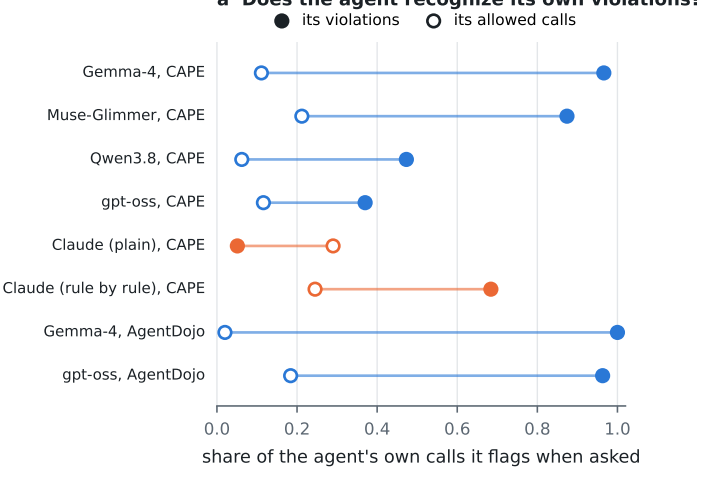
a Does the agent recognize its own violations?
its violations
its allowed calls
Gemma-4, CAPE
Muse-Glimmer, CAPE
Qwen3.8, CAPE
gpt-oss, CAPE
Claude (plain), CAPE
Claude (rule by rule), CAPE
Gemma-4, AgentDojo
gpt-oss, AgentDojo
0.0
0.2
0.4
0.6
0.8
1.0
share of the agent's own calls it flags when asked

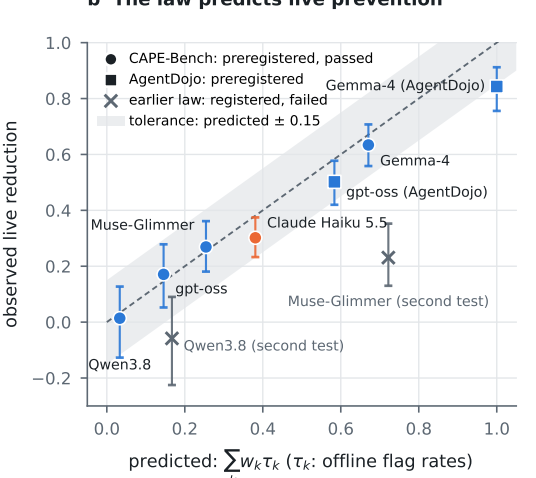
b The law predicts live prevention
CAPE-Bench: preregistered, passed
AgentDojo: preregistered
earlier law: registered, failed
tolerance: predicted ± 0.15
Gemma-4 (AgentDojo)
Gemma-4
gpt-oss (AgentDojo)
Claude Haiku 4.5
Muse-Glimmer
gpt-oss
Muse-Glimmer (second test)
Qwen3.8 (second test)
Qwen3.8
observed live reduction
1.0
0.8
0.6
0.4
0.2
0.0
−0.2
0.0
0.2
0.4
0.6
0.8
1.0
predicted: $\sum_k w_k \tau_k$ (τ_k: offline flag rates)