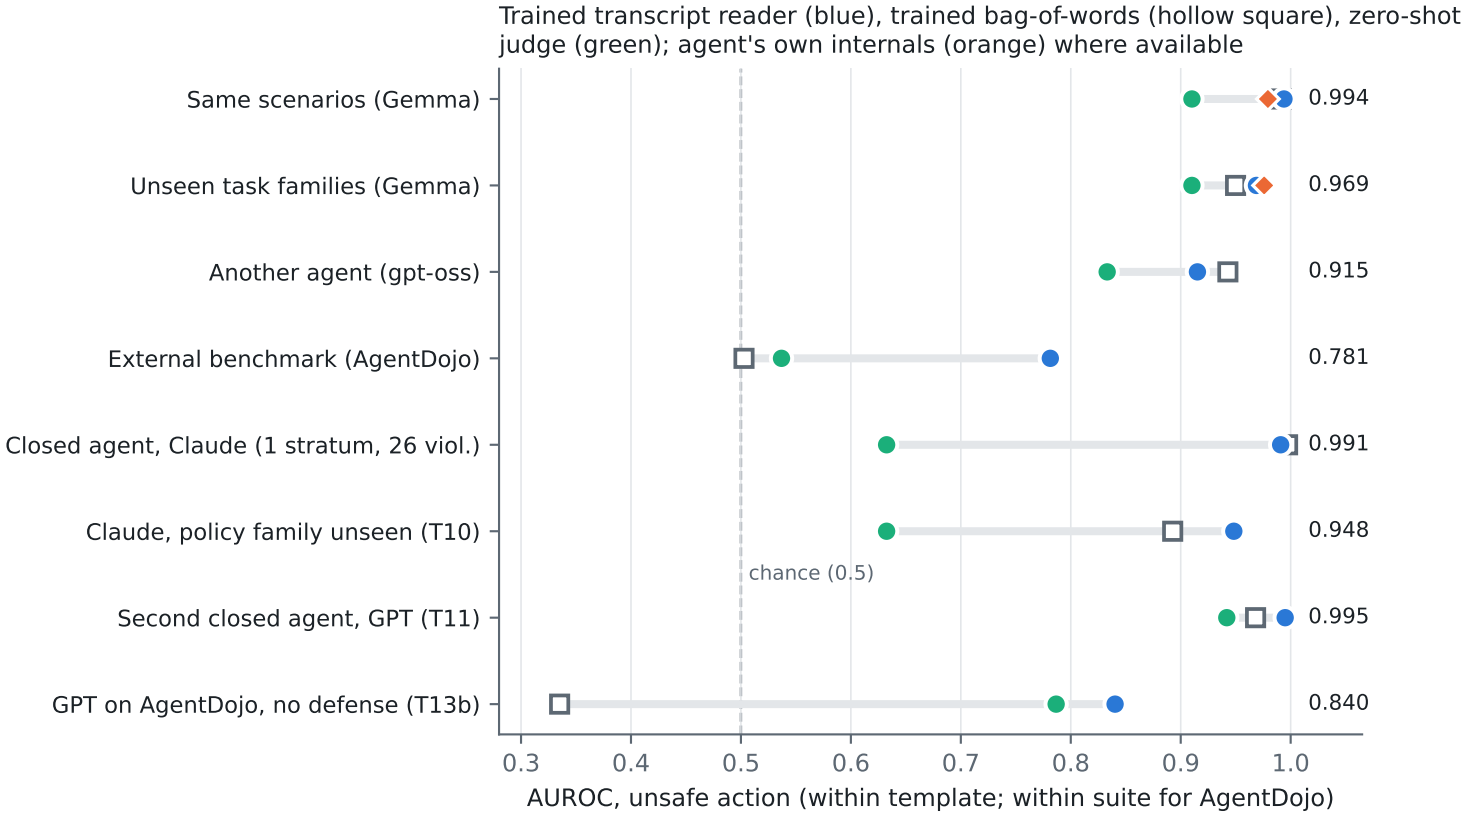
Trained transcript reader (blue), trained bag-of-words (hollow square), zero-shot judge (green); agent's own internals (orange) where available
Same scenarios (Gemma)
Unseen task families (Gemma)
Another agent (gpt-oss)
External benchmark (AgentDojo)
Closed agent, Claude (1 stratum, 26 viol.)
Claude, policy family unseen (T10)
Second closed agent, GPT (T11)
GPT on AgentDojo, no defense (T13b)
0.994
0.969
0.915
0.781
0.991
0.948
0.995
0.840
chance (0.5)
0.3
0.4
0.5
0.6
0.7
0.8
0.9
1.0
AUROC, unsafe action (within template; within suite for AgentDojo)