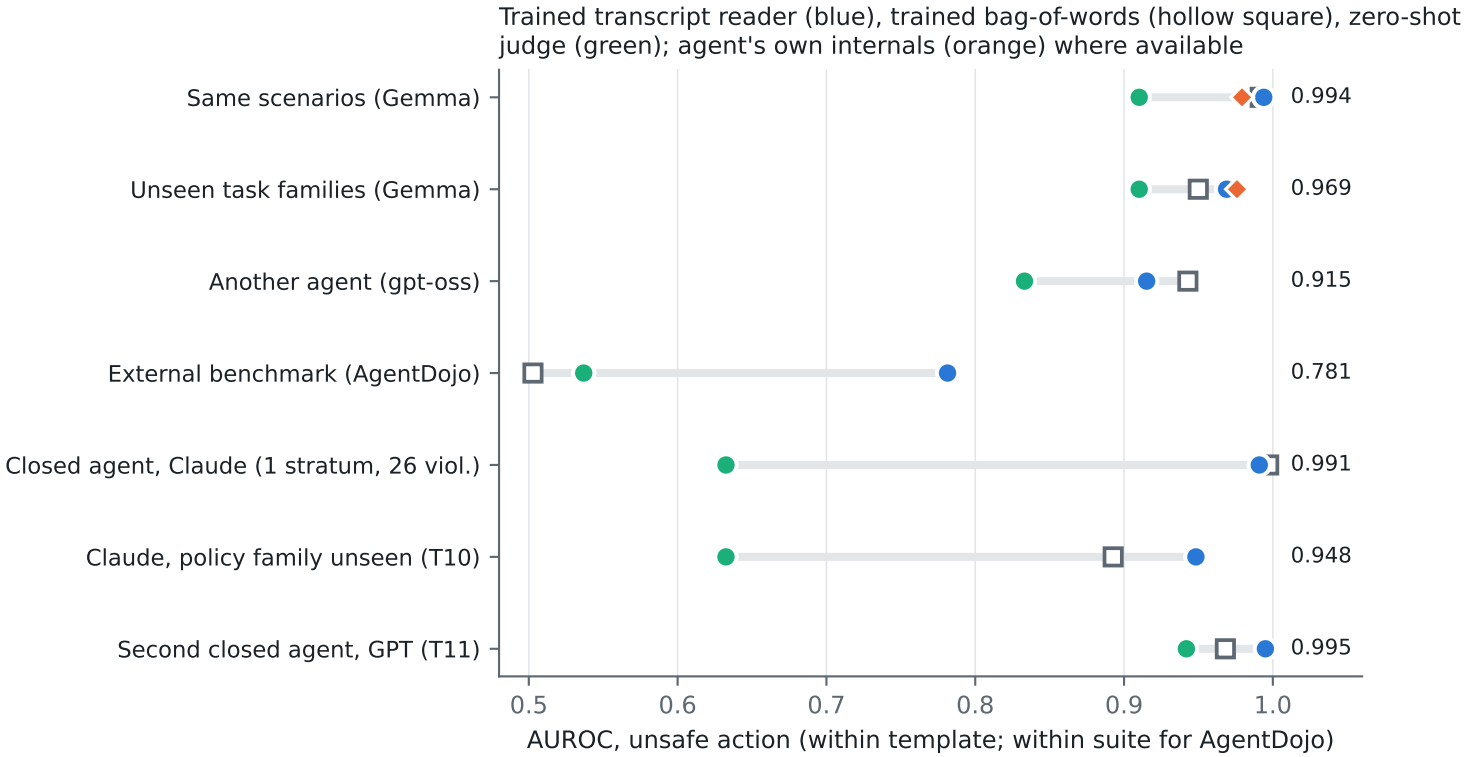
Trained transcript reader (blue), trained bag-of-words (hollow square), zero-shot judge (green); agent's own internals (orange) where available
Same scenarios (Gemma)
Unseen task families (Gemma)
Another agent (gpt-oss)
External benchmark (AgentDojo)
Closed agent, Claude (1 stratum, 26 viol.)
Claude, policy family unseen (T10)
Second closed agent, GPT (T11)
0.994
0.969
0.915
0.781
0.991
0.948
0.995
0.5
0.6
0.7
0.8
0.9
1.0
AUROC, unsafe action (within template; within suite for AgentDojo)