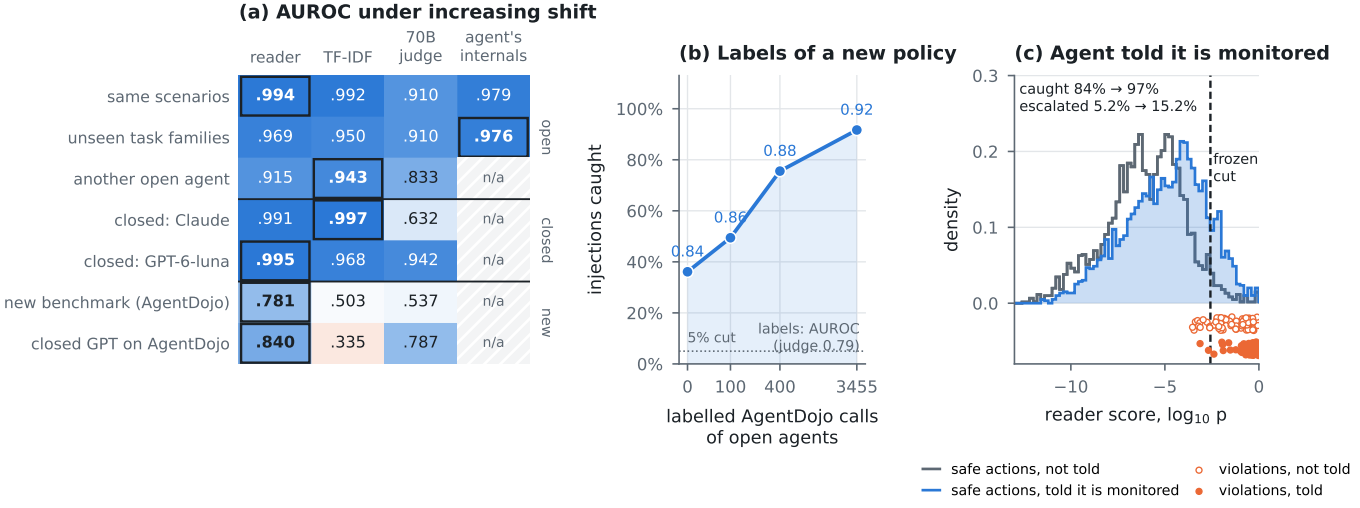
(a) AUROC under increasing shift
reader
TF-IDF
70B judge
agent's internals
same scenarios
.994
.992
.910
.979
unseen task families
.969
.950
.910
.976
another open agent
.915
.943
.833
n/a
closed: Claude
.991
.997
.632
n/a
closed: GPT-6-luna
.995
.968
.942
n/a
new benchmark (AgentDojo)
.781
.503
.537
n/a
closed GPT on AgentDojo
.840
.335
.787
n/a
open
closed
new
(b) Labels of a new policy
injections caught
100%
80%
60%
40%
20%
0%
0.84
0.86
0.88
0.92
5% cut
labels: AUROC (judge 0.79)
0
100
400
3455
labelled AgentDojo calls of open agents
(c) Agent told it is monitored
caught 84% → 97%
escalated 5.2% → 15.2%
frozen cut
density
0.3
0.2
0.1
0.0
−10
−5
0
reader score, log10 p
safe actions, not told
safe actions, told it is monitored
violations, not told
violations, told